

Agri-Climate Risk & Global Environmental Analytics Platform

Complete Project Documentation

Live dashboard	https://agriclimate.streamlit.app/
Repository	https://github.com/lovevyas/Agri_Climate_Data_Analysis_Platform
Cloud	AWS (ap-south-1), provisioned with Terraform
Stack	PySpark on AWS Glue · Lambda · S3 · RDS PostgreSQL · Streamlit

Every number in this document was computed from the pipeline's own output, not estimated. The script that produced them is described in §9.

Table of contents

- 1. [What the project does](#)
 - 2. [Prerequisites](#)
 - 3. [Architecture](#)
 - 4. [Code flow](#)
 - 5. [Interfaces and contracts](#)
 - 6. [Data model](#)
 - 7. [Use cases, code and results](#)
 - 8. [Running the project](#)
 - 9. [Verification](#)
 - 10. [Cost and teardown](#)
-

1. What the project does

Three unrelated CSV sources are combined into one queryable star schema:

Source	Rows	What it holds
agriculture_crop_analysis.csv	2,000	Farm-level crop cycles: sowing, yield, cost, revenue
imd_weather_station_data.csv	2,000	Official India Meteorological Department station readings
noaa_climate_environmental_data.csv	2,000	Global NOAA climate and air-quality observations

The hard part is that farms and weather stations have **no shared key**. They are joined geographically, with a temporal fallback — see [§6.2](#).

The result is served two ways: Parquet files in S3 for analytical tools, and a PostgreSQL star schema for SQL queries and the dashboard.

2. Prerequisites

2.1 To run locally (no AWS account required)

Requirement	Version	Why
Python	3.9+	3.12 used in development
JDK	11 or 17	PySpark runs on the JVM
pip packages	see <code>requirements.txt</code>	dashboard and analysis
PySpark	3.5.1	local ETL only — not in <code>requirements.txt</code>

PySpark is deliberately excluded from `requirements.txt`: AWS Glue ships its own Spark runtime, and including it would add ~300 MB to the Streamlit Cloud build for no benefit.

```
python -m venv .venv
.venv\Scripts\activate      # macOS/Linux: source .venv/bin/activate
pip install -r requirements.txt
pip install pyspark==3.5.1  # only if running the local pipeline
```

Verify the JVM is visible to Spark:

```
java -version
```

2.2 To deploy to AWS

Requirement	Notes
AWS account	Free tier covers <code>db.t3.micro</code> RDS for 12 months
AWS CLI	configured via <code>aws configure</code>
Terraform	1.5+
IAM permissions	create S3, IAM roles, RDS, Glue, Lambda


Confirm credentials resolve before running Terraform:

```
aws sts get-caller-identity
```

2.3 Python dependencies

```
streamlit>=1.30.0      # dashboard framework
pandas>=2.0.0          # dataframes
numpy>=1.24.0          # correlation matrix masking
matplotlib>=3.7.0      # figure rendering
seaborn>=0.12.0        # statistical charts
sqlalchemy>=2.0.0      # RDS connection for the live mode
psycopg2-binary>=2.9.0 # PostgreSQL driver
pyarrow>=14.0.0        # Parquet reader
```


3. Architecture



Design note. Provisioning is entirely declarative (Terraform). `boto3` appears in exactly one place — inside the Lambda — where calling an AWS API at runtime is the actual requirement. Infrastructure is never created imperatively from a script.

4. Code flow

4.1 Module map



4.2 Execution order inside the Glue job

The entry point reads top-to-bottom as Extract → Transform → Load:

```
# aws/glue/glue_job.py

# Extract
raw = {
    name: spark.read.csv(f"{raw_prefix}/{filename}", header=True, schema=schema())
    for name, (filename, schema) in RAW_FILES.items()
}

# Transform
ag_clean = glue_transforms.clean_agriculture(raw["agriculture"])
imd_clean = glue_transforms.clean_imd(raw["imd"])
noaa_clean = glue_transforms.clean_noaa(raw["noaa"])

dim_farms, dim_weather_stations = glue_transforms.build_dimensions(ag_clean, imd_clean)
matched_farms = glue_transforms.match_farms_to_stations(ag_clean, imd_clean,
```

```

dim_weather_stations)

    tables = {
        "dim_farms": dim_farms,
        "dim_weather_stations": dim_weather_stations,
        "fact_crop_yield": glue_transforms.build_fact_crop_yield(matched_farms),
        "fact_weather_observations":
glue_transforms.build_fact_weather_observations(imd_clean),
        "fact_global_climate_environment": noaa_clean,
    }

    # Load
    glue_load.write_curated_parquet(tables, curated_prefix)

    url, props = glue_load.jdbc_settings(args)
    glue_load.reset_tables(args)
    glue_load.load_star_schema(tables, url, props)

    summaries = glue_analytics.build_all(...)
    glue_analytics.load_all(summaries, url, props)

```

4.3 Why the ETL is split across five files

Glue executes **one** entry-point script. Splitting it means the other four modules must be shipped separately, which Terraform handles by zipping them:

```

# infra/main.tf
locals {
    glue_modules = ["glue_schemas.py", "glue_transforms.py", "glue_load.py", "glue_analytics.py"]
}

data "archive_file" "glue_modules" {
    type      = "zip"
    output_path = "${path.module}/../aws/glue/glue_modules.zip"

    dynamic "source" {
        for_each = local.glue_modules
        content {
            content = file("${path.module}/../aws/glue/${source.value}")
            filename = source.value
        }
    }
}

```

```

# infra/glue.tf – the zip is handed to Glue here
default_arguments = {
    "--extra-py-files" =
"s3://${aws_s3_bucket.data_lake.id}/${aws_s3_object.glue_modules.key}"
    "--additional-python-modules" = "psycopg2-binary"
}

```

Maintenance warning. Adding a new module under `aws/glue/` requires adding its filename to `local.glue_modules`. Forgetting this fails at runtime with `ModuleNotFoundError`, not at `terraform`

```
apply .
```

```
resource "aws_s3_bucket_notification" "csv_upload" {
  bucket = aws_s3_bucket.data_lake.id

  lambda_function {
    lambda_function_arn = aws_lambda_function.trigger_glue.arn
    events               = ["s3:ObjectCreated:*"]
    filter_prefix        = "raw/"
    filter_suffix        = ".csv"
  }

  depends_on = [aws_lambda_permission.allow_s3]
}
```

[illegible]

```
except Exception as e:
    return {"statusCode": 500, "body": json.dumps({"error": str(e)})}
```

Response contract

Status	Body
200	{"message", "file", "job_run_id"}
500	{"error"}

5.3 Glue job arguments

Seven arguments, all required. Missing any one aborts the run with `GlueArgumentError` :

Argument	Source
--JOB_NAME	supplied by Glue
--S3_BUCKET_NAME	Terraform output <code>bucket_name</code>
--RDS_HOST	Terraform output <code>db_endpoint</code> (host portion)
--RDS_PORT	5432
--RDS_DB	<code>agri_climate</code>
--RDS_USER	<code>dbadmin</code>
--RDS_PASSWORD	Terraform output <code>db_password</code> (sensitive)

5.4 Output targets

Target	Location	Partitioned by
<code>dim_farms</code>	<code>s3://bucket/curated/dim_farms</code>	—
<code>dim_weather_stations</code>	<code>s3://bucket/curated/dim_weather_stations</code>	—
<code>fact_crop_yield</code>	<code>s3://bucket/curated/fact_crop_yield</code>	<code>season</code>
<code>fact_weather_observations</code>	<code>s3://bucket/curated/fact_weather_observations</code>	<code>season</code>
<code>fact_global_climate_environment</code>	<code>s3://bucket/curated/...</code>	<code>country</code>

All eight tables are additionally written to RDS over JDBC.

5.5 Dashboard navigation

The dashboard's "routes" are a sidebar radio mapped to view modules:

```
# visualizations/app.py
PAGES = {
    "Platform Overview": overview,
    "Use Case 1: Yield & Weather": usecase1_yield_weather,
    "Use Case 2: Climate Risk": usecase2_climate_risk,
```



```

    "Use Case 3: NOAA Environment": usecase3_noaa_environment,
    "Hidden Insights": hidden_insights,
}

PAGES[page].render(dfs)

```

The directory is named `views/`, not `pages/`. Streamlit treats a `pages/` folder beside the entry point as automatic multi-page navigation, which would have conflicted with this manual router.

6. Data model

6.1 Star schema




Table	Rows	Columns
dim_farms	2,000	10
dim_weather_stations	2,000	4
fact_crop_yield	2,000	25
fact_weather_observations	2,000	21
fact_global_climate_environment	2,000	31
analytics.yield_weather_correlation	rollup	6
analytics.climate_risk_farms	rollup	7
analytics.global_environment_summary	rollup	7

6.2 Farm-to-station matching

Farms and stations share no key. Two strategies, in order:

```

# aws/glue/glue_transforms.py
def match_farms_to_stations(ag_clean, imd_clean, dim_weather_stations):
    keep_cols = ag_clean.columns + ["matched_station_code", "match_type"]

    exact_raw = (
        ag_clean.join(dim_weather_stations, on=["state", "district"], how="inner")

```

```

        .withColumn("matched_station_code", col("station_code"))
        .withColumn("match_type", lit("Exact District Match"))
    )
    # A district can hold several stations, which would fan one farm out into
    # several rows and collide with fact_crop_yield's farm_id primary key.
    exact_matches = (
        exact_raw
        .withColumn("_rn",
row_number().over(Window.partitionBy("farm_id").orderBy("station_code")))
        .filter(col("_rn") == 1)
        .select(*keep_cols)
    )

    unmatched = ag_clean.join(dim_weather_stations, on=["state", "district"], how="left_anti")

    # Project to just the join keys: carrying all of imd_clean here would
    # duplicate column names (season, rainfall_mm...) and make later
    # references ambiguous.
    station_readings = imd_clean.select("station_code", "state", "observation_date")
    fallback = (
        unmatched.join(station_readings, on="state", how="inner")
        .withColumn("date_diff", spark_abs(datediff(col("sowing_date"),
col("observation_date"))))
        .withColumn("_rn", row_number().over(
            Window.partitionBy("farm_id").orderBy("date_diff", "station_code")
        ))
        .filter(col("_rn") == 1)
        .withColumn("matched_station_code", col("station_code"))
        .withColumn("match_type", lit("Nearest In State"))
        .select(*keep_cols)
    )

    return exact_matches.unionByName(fallback)

```

Result

Match type	Farms	Share
Exact District Match	1,739	87.0%
Nearest In State	261	13.0%
Total	2,000	100%

182 distinct stations are actually used out of 2,000 available.

6.3 Loading strategy

Spark's `mode="overwrite"` issues `DROP TABLE`. That fails against foreign keys and would replace the DDL's `DECIMAL` types with Spark's inferred ones. Instead:

```

# aws/glue/glue_load.py
RESET_SQL = """
    TRUNCATE TABLE
        fact_crop_yield, fact_weather_observations, fact_global_climate_environment,
        dim_farms, dim_weather_stations,

```

```

        analytics.yield_weather_correlation, analytics.climate_risk_farms,
        analytics.global_environment_summary
    RESTART IDENTITY CASCADE;
"""

def load_star_schema(tables, url, props):
    """Load dimensions before facts so the foreign keys resolve."""
    print("Loading Dimension Tables to RDS...")
    for name in ("dim_farms", "dim_weather_stations"):
        tables[name].write.jdbc(url=url, table=name, mode="append", properties=props)

    print("Loading Fact Tables to RDS...")
    for name in ("fact_crop_yield", "fact_weather_observations",
"fact_global_climate_environment"):
        tables[name].write.jdbc(url=url, table=name, mode="append", properties=props)

```

Postgres only permits truncating an FK-referenced table if its dependents are truncated in the **same** statement — hence one command for all eight tables.

7. Use cases, code and results


7.1 Reading these results honestly

The source data is synthetic. Some specified charts plot variables that are statistically independent, and the documentation says so rather than inventing a narrative. Each such chart is paired with a companion chart built on a relationship that does hold.

Use Case 1 — Yield vs. official weather


Question. Does the weather recorded at official stations explain farm profitability?

Chart 1.1 — Profitability vs. station rainfall



Flat cloud, no left-to-right slope, nearly all of it below break-even. $r = +0.048$ — station rainfall does not separate profitable farms from unprofitable ones.

Chart 1.2 — Rainfall reporting gap by state



The measurable disagreement between farmer-reported and official rainfall. This is the genuinely useful output of the join: it quantifies data-quality drift, which matters for insurance claims assessed on self-

reported weather.

Chart 1.3 — Profit driver explorer (interactive)


```
# visualizations/views/usecase1_yield_weather.py
PROFIT_DRIVERS = {
    "Official Station Rainfall (mm)": "rainfall_mm",
    "Production Cost (INR)": "production_cost_inr",
    "Land Area (acres)": "land_area_acres",
}

def _render_profit_explorer(merged):
    st.write("#### 3. Profit Driver Explorer (interactive)")

    all_crops = sorted(merged["crop"].unique())
    col_a, col_b = st.columns([1, 1])
    with col_a:
        crops = st.multiselect("Crops to compare", all_crops, default=all_crops[:4])
    with col_b:
        driver = st.radio("Compare profit against", list(PROFIT_DRIVERS), index=0)

    fig = graphs_insights.plot_profit_trend(merged, PROFIT_DRIVERS[driver], driver, crops)
    st.pyplot(fig)
```

The chart binds x-axis to a selector so the user discovers the answer by comparison:



Quantile bins are used rather than equal-width so each point rests on a similar number of farms:

```
# visualizations/graphs_insights.py
def plot_profit_trend(merged_df, x_col, x_label, crops):
    """Mean profit across quantile bins of x_col, one line per selected crop."""
    fig, ax = plt.subplots(figsize=(10, 5))

    for crop in crops:
        sub = merged_df[merged_df["crop"] == crop]
        if sub.empty:
            continue
        # Quantile bins keep each point backed by a similar number of farms,
        # so a line never swings wildly on the strength of one outlier.
```

```

binned = pd.qcut(sub[x_col], q=5, duplicates="drop")
summary = sub.groupby(binned, observed=True)["profit_loss_inr"].mean()
ax.plot([b.mid for b in summary.index], summary.values, marker="o", label=crop)

ax.axhline(0, color="red", linestyle="--", alpha=0.5, label="Break-even")
return fig


```

Finding. `profit_loss_inr` equals `market_price_inr - production_cost_inr` **exactly** (max residual `1.16e-10`). Production cost correlates with profit at $r = -0.887$; station rainfall at $r = +0.048$. Cost runs roughly twice revenue throughout, which is why 95.4% of farms show a loss and the median margin is -114% . Profitability here is a cost-structure story, not a climate one.

Use Case 2 — Climate risk and farm resilience


Question. Do official weather alerts predict crop damage?

Chart 2.1 — State × crop health




Scores cluster in the 60s–70s; no state-crop pairing stands out dramatically.

Chart 2.2 — Pest rate under alert vs. no alert (*specified chart*)



33.9% vs 32.2% — a 1.7-point difference. Flood and cyclone alerts do **not** predict pest attacks. Weather alerts are not a usable early-warning signal here.

Chart 2.3 — What a pest attack actually costs (*companion chart*)



```
# visualizations/views/usecase2_climate_risk.py
pest_summary = (
    yield_df.groupby("pest_attack")[["crop_health_score",
    "yield_gap_tonnes"]].mean().reset_index()
)
fig = graphs_insights.plot_pest_impact(pest_summary)
st.pyplot(fig)
```


Metric	No pest	Pest attack	Change
Crop health score	76.66	45.10	−31.6 points
Yield gap (tonnes)	−1.21	+10.26	11.5 t swing
Actual yield (tonnes)	50.63	41.05	−19%

Finding. Pests are the single largest driver of lost yield (`crop_health_score` vs `yield_gap_tonnes` , $r = -0.616$) — but weather alerts give no warning of them. Resource pre-positioning should therefore be driven by pest surveillance, not meteorological alerts.

Use Case 3 — Global environmental monitoring


Question. What drives air quality and disaster risk?

Chart 3.1 — AQI by country (*specified chart*)




Canada 106.7 · Mexico 107.0 · USA 106.2 — within one point. Country carries no information about air quality in this dataset.

Chart 3.2 — Climate zone × disaster risk (*specified chart*)



Every cell near 33%. Climate zone carries no information about disaster risk.

Chart 3.3 — Environmental correlation matrix (*companion chart*)



```
# visualizations/graphs_insights.py
def plot_env_correlation_heatmap(noaa_df):
    """Correlation matrix of the NOAA environmental measures."""
    cols = [
```

```


    "avg_temperature_c", "evapotranspiration_mm", "precipitation_mm",
    "humidity_percent", "soil_moisture_percent", "air_quality_index", "pm25", "pm10",
]
corr = noaa_df[cols].corr()

fig, ax = plt.subplots(figsize=(10, 6))
sns.heatmap(
    corr,
    mask=np.triu(np.ones_like(corr, dtype=bool)), # upper half mirrors the lower, so hide
    annot=True, fmt=".2f", cmap="RdBu_r", center=0, vmin=-1, vmax=1,
    linewidths=0.5, cbar_kws={"label": "Correlation (r)"}, ax=ax,
)
return fig

```

Same chart type as 3.2, real structure. Two clusters emerge: a thermal one (temperature → evapotranspiration, **+0.945**) and a hydrological one (precipitation → humidity → soil moisture, and precipitation → AQI at **-0.741**).

Chart 3.4 — Drought index profile (*companion chart*)




Drought index	Soil moisture	Precipitation
Low	43.93%	30.54 mm
Medium	36.03%	11.65 mm
High	32.29%	3.26 mm

Finding. Unlike climate zone, the drought index is a genuine summary of local water availability — both measures fall monotonically as severity rises.


Hidden Insights

Evapotranspiration vs. temperature




0.8 mm below 0 °C rising to 6.0 mm above 30 °C — roughly sevenfold ($r = +0.945$). A direct visual for water stress: higher temperatures mean faster soil drying.

Actual vs. expected yield



$r = +0.972$. Farms forecast well; points below the red line are the shortfalls, averaging about 5% of expected yield.

AQI vs. precipitation



$r = -0.741$. Rain scrubs particulates from the air. This is the effect the flat AQI-by-country chart was missing.

8. Running the project

8.1 Locally

```
python pyspark/run_pipeline.py      # raw CSVs -> datasets/curated/ Parquet
python testing/verify_tests.py      # writes testing/test_results_log.csv
streamlit run visualizations/app.py
```

8.2 Deploying to AWS

```
cd infra
terraform init
terraform apply
```

Create the tables once (reads `DB_HOST` and `DB_PASSWORD` from the environment):

```
python apply_schema.py
```

From here the pipeline is event-driven:

```
aws s3 cp ../datasets/raw/agriculture_crop_analysis.csv s3://<bucket>/raw/
```

To start a run directly instead:

```
.\test_glue_args.ps1
```

PowerShell note. `ConvertTo-Json | Set-Content -Encoding utf8` writes a UTF-8 BOM in PowerShell 5.1, which the AWS CLI rejects when reading `file://`. The helper script uses `[System.IO.File]::WriteAllText` with `UTF8Encoding($false)` instead.

8.3 Deployed dashboard

Hosted on Streamlit Community Cloud from `visualizations/app.py`, reading the committed Parquet files. The live-RDS option is intentionally unused there — connecting would require typing a database password into a public page.

9. Verification

The pipeline was validated at four levels:

Level	Check	Result
Schema	Fact column lists vs <code>create_tables.sql</code>	25 and 21 columns, exact match
Referential	<code>matched_station_code</code> values absent from <code>dim_weather_stations</code>	0 orphans
Key integrity	Duplicate PKs in each fact table	none
Runtime	Glue job end-to-end on AWS	succeeded
Dashboard	All 5 pages, 13 charts	0 exceptions
Dependencies	Every import present in <code>requirements.txt</code>	all covered

9.1 Bugs found and fixed during development

Symptom	Root cause	Fix
Reference 'season' is ambiguous	Fallback join carried all of <code>imd_clean</code> , duplicating column names	Project to the three join keys only
cannot drop table <code>dim_farms</code>	Spark <code>mode="overwrite"</code> issues <code>DROP TABLE</code> , blocked by FKs	<code>TRUNCATE ... CASCADE + mode="append"</code>
Column "state" not found	Facts carried dimension columns; <code>append</code> validates against the real table	Explicit <code>.select()</code> to the DDL columns
duplicate key ... <code>fact_crop_yield_pkey</code>	A district holds several stations, fanning one farm into several rows	Window dedup to one station per farm
Overview showed 2,000 "stations matched"	Metric counted the dimension, not actual matches	Count <code>matched_station_code</code> distinct → 182
5 charts would break on seaborn 0.14	<code>palette=</code> without <code>hue=</code> is deprecated	Added <code>hue=</code> + <code>legend=False</code>

The `mode="overwrite"` bug is worth dwelling on: it was silently *masking* the schema error. Overwrite recreated each table to match whatever the DataFrame held, so the extra dimension columns never raised anything — they just quietly corrupted the star schema. Switching to `append` made Spark validate against the real table and surface the mismatch.

10. Cost and teardown

Service	Billing	Notes
RDS <code>db.t3.micro</code>	hourly	free tier 12 months; the only continuous cost
S3	per GB	~2 MB stored — negligible
Glue	per DPU-hour	2 × G.1X, a few minutes per run
Lambda	per invocation	free tier covers it

RDS is the only meaningful cost. When finished:

```
cd infra
terraform destroy
```

Security note

The RDS instance is deliberately `publicly_accessible = true` with ingress from `0.0.0.0/0`, so the dashboard can reach it from a laptop without a bastion host or VPN. It is protected only by a generated 20-character password.

This is a demo convenience, not a pattern to copy. Production would place RDS in a private subnet and reach it through a bastion, a VPN, or a VPC-attached Lambda. The Terraform state holding that password is excluded from version control.